When Courts Miscite: A Provable Baseline for Citation Errors in Swiss Judicial Decisions

Jonas Hertner[1]

[1]OpenCaseLaw

August 14, 2026

Abstract

Measurement studies report legal hallucination in 58–88 % of responses to case-law queries put to large language models [3] and in 17–33 % for commercial legal-RAG tools [8]. These figures lack a reference point: how often do *courts themselves* cite authority that does not exist? We give, to our knowledge, the first decidability-based answer for a national jurisdiction. The Swiss Federal Supreme Court's official reporter (BGE/ATF/DTF) is a closed, continuously paginated series, held completely in our corpus from 1875; the nonexistence of a cited locus is therefore decidable. Scanning 584,524 distinct reporter-citation edges (one per decision and cited locus) in 119,214 Swiss decisions issued since January 2024, we find 591 provably nonexistent citations — 1,011 per million (one in 989), Wilson 95 % 933–1,096 ppm, decision-cluster bootstrap 931–1,093 ppm. The machine studies measure responses and count existent-but-wrong citations; we measure citations and count provable nonexistence only. The units are not divisible into a ratio, so we state them side by side: roughly one citation in a thousand in the published human record, against 17–88 % of machine responses under broader error definitions. The court's public resolver independently confirms the scan — 542 of 591 findings carry a token the resolver 404s, none resolves — and because the series is closed, most errors carry a local-damage signature: for 272 findings the cited page exists under a sibling division of the same volume, and where the citing document itself states the intended citation, the deterministic repair reproduces it exactly. Errors concentrate in reused boilerplate — one mistyped division letter recurs verbatim in three decisions across two court levels. Reported to the issuing courts on 3 August 2026, the findings drew replies from seven court administrations within three days; two courts confirmed corrections, one already implemented. A record that can be audited can be healed. Every scan-derived number regenerates from released, hash-manifested artifacts.

1 Introduction

The reliability of legal citations has become a measured quantity. Dahl et al. [3] report hallucination in 58–88 % of responses to reference-based case-law queries put to general-purpose LLMs; Magesh et al. [8] measure 17–33 % for commercial legal-RAG products; a public registry documents over 1,300 court proceedings in which judges confronted AI-generated citations in filings [2]. Detection work has followed: verification against a national citation graph [10], benchmarks with injected hallucinations [7], and surveys of detection tools [1]. Verifying citations is now framed as a burden courts themselves need help with [11].

Every one of these numbers floats without a baseline. Fabrication rates are alarming relative to an implicit assumption — that human legal writing cites what exists — which no study quantifies. The nearest observation is an aside: a hallucination-detection benchmark built on pre-LLM briefs notes "over 10 human-made citation typos" in its source material and moves on [7]. Measuring the human side properly requires something unusual: a corpus in which the *nonexistence* of a cited authority is provable rather than merely unresolved. Reference-error studies in other fields illustrate the difficulty: quotation-error rates near 20 % are reported in medical literature [9], but against open-ended reference universes where absence of evidence remains evidence of nothing.

Swiss law provides the unusual conditions. The Federal Supreme Court's official reporter — *BGE* in German, *ATF* in French, *DTF* in Italian; one series, three prefixes — is closed and continuously paginated: volume and division fix a page range, decisions begin at known pages, and the series is complete in our corpus from 1875 to the present [5]. For a citation of the form *BGE 148 I 356*, existence is therefore decidable: either some reported decision begins at (or plausibly contains) that locus, or none does. "Provably nonexistent" is a statement about a finite, enumerable set, not a retrieval failure.

We scan the reporter citations of every Swiss decision issued since January 2024 in the corpus — 584,524 distinct citation edges (the extractor counts each cited locus once per decision) across 119,214 decisions from federal and cantonal courts in all three official languages — and prove nonexistence for 591: 1,011 per million (one in 989). The baseline the field has been missing can then be stated next to the machine numbers, each in its own unit: roughly 0.1 % of citations in the published human record, against 17–33 % of commercial legal-RAG responses and 58–88 % of raw-LLM responses under broader error definitions — the differences in unit, universe and definition are laid out in Table 3 rather than divided away.

Three further contributions come from reading the errors rather than counting them. First, a **mechanism taxonomy**, computed deterministically over every finding rather than curated from examples: human miscitation overwhelmingly matches local-damage signatures of a true citation — a division letter substituted while the page stays right, a digit doubled or dropped, a year where a page belongs — where LLM fabrication invents whole plausible loci. In 272 of 591 findings the cited page exists as a decision start under a *sibling division of the same volume*. Second, **propagation**: identical erroneous tokens recur across decisions, courts and years, carried by reused doctrinal passages — the token *BGE 147 IIII 249*, division III with a doubled letter, appears verbatim in three decisions across two court levels, each quoting the same standard. Third, a **remediation loop with observed outcomes**: the findings were verified, reported to the issuing courts on 3 August 2026, and drew replies from seven court administrations within three days, including two confirmed corrections (Section 4.4). A record that can be audited can also be healed.

What this paper is not. It does not rank courts, and it does not claim the found errors affected outcomes: a wrong volume number beside a correct parallel citation is a blemish, not a miscarriage. The rate we report is a property of the published record, offered as the denominator the machine-fabrication literature lacks.

2 Related work

LLM citation fabrication. Dahl et al. [3] established hallucination rates for general-purpose models over fourteen tasks on the U.S. federal judiciary — nine reference-based, five reference-free; their pooled 58–88 % covers the reference-based tasks, which span existence, court, citation, author, disposition and more; Magesh et al. [8] extended measurement to commercial legal-RAG tools with 202 preregistered queries; Charlotin [2] tracks court-confronted instances. Detection approaches verify model output against citation graphs [10], benchmark detectors on synthetic corruptions [7], or survey tool capabilities [1]. All measure or police the machine side; none that we found measures the human side, which is this paper's object.[1]

[1] Sweep documented in the released `notes/novelty-audit.md`: the LLM-hallucination literature and its surveys, legal citation-infrastructure studies, and reference-error studies in scientific writing. Bluebook-error, citator-accuracy and non-English empirical literatures are searched but not exhaustively; the claim is scoped "to our

Reference errors in human writing. Citation and quotation error rates are long studied in scientific literature [9], with majorities of surveyed articles containing at least minor reference defects. Those studies rely on manual lookup in open-ended bibliographic universes; the closed-world reporter series lets us replace sampling with exhaustive scanning and judgment with decidability, at the price of covering only reporter-form citations. For judicial opinions specifically, the nearest empirical work measures *link rot*: 29 % of internet links cited in U.S. Supreme Court opinions 1996–2010 no longer worked [6]. Rot is post-publication decay of the referent; our object is publication-time nonexistence of the cited authority.

Swiss legal data. The corpus and citation graph underlying this study are described in the companion resource paper [5]; the Swiss Federal Supreme Court Dataset [4] provides structured metadata for the federal layer.

3 Method

3.1 Decidability conditions

A reporter citation names a volume v, a division d, and a page p. The series is continuously paginated within each (volume, division): decisions begin at strictly increasing pages, and a decision occupies the pages from its start to the next decision's start. Our corpus holds every reported decision since volume 1 (1875), giving for each (v, d) the complete sorted list of start pages $S_{v,d}$ (Figure 1).

A cited locus (v, d, p) is **provably nonexistent** iff one of:

1. *Division absent*: $S_{v,d} = \emptyset$ although volume v exists (the reporter's divisions are I, Ia, Ib, II, III, IV, V; the Ia/Ib split existed 1955–1994 and is evaluated as a family with I, since citations casefold the sub-letter).
2. *Page beyond the series*: $p > \max S_{v,d} + w_{v,d}$, where the window $w_{v,d}$ is the longest observed decision length in that (v, d), with a floor of 30 pages — a mid-volume page can always be a legitimate deep pin-cite, so only pages beyond the last decision's plausible extent are provable.
3. *Volume out of range*: v exceeds the newest volume by more than one (the newest two volumes are still filling and are exempt from all three conditions).

Everything else — including citations that merely fail to resolve — is *not* counted, and nonexistence is asserted as of the citing decision's date (a volume beyond next year's does not exist when cited, though it may later).

The three conditions differ in epistemic strength. Division-absent and volume-out-of-range are decidable from the series structure alone. Page-beyond-the-series depends on the window $w_{v,d}$, which is inductive — the final decision of a (v, d) could in principle exceed every predecessor's length. We therefore close the gap externally: the court's own resolver redirects any page *inside* a decision to that decision and returns 404 otherwise (Section 3.2), so a 404 is an authoritative statement that no decision contains the cited page. Of the 591 findings, 542 carry a token the resolver 404s; the remaining 49 cite volumes below the resolver's coverage floor and rest on the corpus proof and the window rule alone.

3.2 Guards against false positives

Five mechanisms, each added after a concrete failure it now prevents:

knowledge" accordingly.

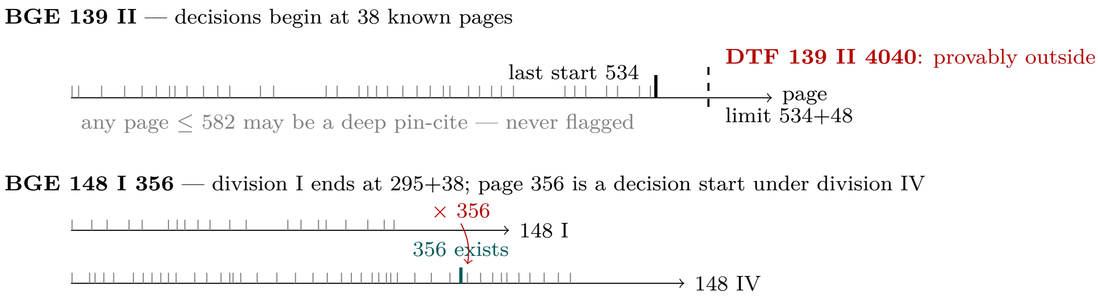

Figure 1: Decidability on the closed series, drawn from the released index. Top: continuous pagination fixes, for every (volume, division), the last decision start and an adaptive window; only pages beyond that limit are flagged. Bottom: *BGE 148 I 356* is provably nonexistent — and the same page is a decision start under division IV, the deterministic *division substitution* signature.

- **Own-OCR exclusion.** A corpus-wide scan back to 1875 flags ∼1,500 tokens, concentrated in 19th-century decisions where the flagged token is our OCR of the source, not the court's text. The study window (since 2024) is born-digital, and every finding's context excerpt is carried in the release for inspection.
- **Prefix discipline for bare tokens.** French and Italian typography breaks citations across lines, and older French texts write dates with Roman months (*31 III 2004*). The primary universe is therefore *prefixed* citations only. A token without its own BGE/ATF/DTF prefix is admitted to a separately reported secondary channel only if the source text shows the prefix immediately before it and the preceding character is not a digit (excluding digit-split artefacts like *1 43 V 409*); this channel contributed 1 finding, excluded from every headline number. Of 1,188 raw flags, 596 failed these guards and were dropped.
- **Pin-cite window.** Continuous pagination makes any mid-volume page a possible deep pin-cite; the adaptive window $w_{v,d}$ (Section 3.1) prevents the class of false flags where a long leading case is cited 30+ pages into its body.
- **Open-volume exemption.** The newest volumes are still publishing; nothing in them (or one volume beyond) is decidable and nothing is flagged.
- **Grammar coverage.** The extractor recognises divisions written I–V. Citations using the historical sub-divisions Ia/Ib (volumes 81–120, 1955–1994) — *BGE 120 Ia 45* — are not extracted and enter neither numerator nor denominator: a raw-text count puts this class at 10,723 occurrences in 6,826 window decisions, 1.4 % of the reporter-citation surface. The rate is therefore a property of I–V-form citations (Section 5).
- **External cross-check.** Every distinct flagged token was probed against the Federal Supreme Court's public resolver, whose semantics we verified directly: a decision start returns its own page, an interior page redirects to the containing decision (existence proven), a nonexistent locus returns 404. Of 292 distinct tokens, 38 lie in volumes below the resolver's coverage floor (volume 80, 1954; real older citations also 404 and rest on the corpus proof). All 254 tokens within coverage return 404 and 0 resolve — zero scan false positives on the testable set. The resolver also *positively* confirms the repair candidates of Section 4.2: 109 of 115 unique candidates exist; the remainder are themselves below the coverage floor, none is refuted.

Table 1: Provably nonexistent reporter citations in Swiss decisions issued since 2024-01-01 (scan of 2026-08-13, graph of 2026-08-13). Wilson 95 % intervals; the headline per-token rate additionally carries a cluster-bootstrap interval by decision (10,000 replicates).

Unit	Errors	Universe	Rate (ppm)	95 % CI (ppm)
Citation tokens	591	584,524	1,011	933–1,096 (cluster 931–1,093)
Decisions (≥1 finding)	587	87,063	6,742	6,220–7,308
Distinct cited loci	292	16,708	17,477	15,598–19,577

3.3 Voice

A court's text sometimes *reports* a party's erroneous citation rather than committing one. We screen all findings with a deterministic quotation-marker detector (typographic quotation characters in the context excerpt): 49 of 591 findings (8.3 %) carry markers. Reading each flagged context, three are attributed to a party rather than the court: an appellate court flags the appellant's *BGE 419 III 210* as erroneous and supplies the correction (the citation enters our graph because the erroneous token appears verbatim in the decision text); one decision summarises a debtor's contention in reported speech; one enumerates the authorities the parties cite while holding that none supports them. The sensitivity row of Table 5 excludes all 49 marker-carrying findings and moves the rate by about 8 %. Reported speech without typographic markers remains a residual risk, stated in Section 5; the reading is a single-reader pass by the author, not a protocolised multi-annotator adjudication.

3.4 Denominator

The denominator is counted in the same pass that produces the findings: every distinct prefixed (decision, locus) edge the classifier considered — 584,524 edges from 119,214 decisions, of which 87,063 cite the reporter at least once. The unit is deliberately the edge, not the mention: the extractor counts a cited locus once per decision, so a decision repeating the same wrong citation five times contributes one finding and one denominator entry, and resolved edges carrying several graph target rows are counted once. Per-language and per-form denominators, and the per-decision cluster structure used for the bootstrap, are released with the data.

4 Results

4.1 The rate

Table 1 states the headline three ways. Per citation edge: 591 of 584,524 distinct prefixed reporter citations are provably nonexistent, 1,011 per million (one in 989). Per decision: 0.67 % of the 87,063 decisions that cite the reporter contain at least one provable error. Per distinct locus: 292 distinct erroneous tokens against 16,708 distinct cited loci. The window is a census of the corpus, so the intervals are not sampling corrections for these counts; they quantify uncertainty about the decision-generating process, reading the window as one draw from it. Beside the Wilson interval we report a bootstrap that resamples decisions; the two nearly coincide because the extractor's edge-per-decision unit already removes within-decision repetition (587 decisions carry the 591 findings). Dependence *across* decisions — the shared boilerplate of Section 4.3 — is not captured by either interval; the distinct-locus row is the view least affected by it.

Per language (Table 2): German 915, French 1,081, Italian 970 per million. A homogeneity test does not detect a difference (χ^2=3.8, 2 df) — absence of evidence for a difference, not evidence of equality. Across court classes the data do show a difference: federal decisions run at 877 and cantonal decisions at 1,064 per million (nominal two-sided p=0.04, Table 5); the federal stratum still contributes 147 findings in the federal courts' own published texts.

Table 2: Per-language rates over distinct prefixed citation edges, with Wilson 95 % intervals.

Language	Errors	Tokens	Rate (ppm)	Wilson 95 %
German	210	229,561	915	799–1,047
French	356	329,180	1,081	975–1,200
Italian	25	25,783	970	657–1,431

Table 3: What each study measures. The rates are not directly comparable — the units differ — which is precisely why the human baseline must be stated in its own unit rather than inferred from the machine studies.

	Dahl et al. 2024	Magesh et al. 2025	This study
Unit	response to a case-law query	response to a legal research query	distinct citation edge in a published decision
Population	4 general LLMs; stratified federal case sample	3 commercial legal-RAG tools + GPT-4	every Swiss decision issued since 2024 in the corpus
Universe	14 QA tasks: 9 reference-based, 5 reference-free; n=5,000 per court level, n=100 high-complexity	202 preregistered queries	584,524 prefixed reporter-citation edges
Detection	reference-based: contradiction with known metadata; reference-free: self-contradiction	human coding of responses and cited sources	decidable nonexistence against the complete reporter series, cross-checked on the court's resolver
Error definition	response inconsistent with legal facts (existence, court, citation, author, disposition, . . .)	false statement, or source that does not support the claim	cited locus provably absent from the series
Includes wrong-but-existent citations	yes	yes	no (floor)
Headline	58–88 % of responses, reference-based tasks pooled by model	17–33 % of responses (GPT-4 43 %)	1,011 ppm of citations

Table 3 aligns this measurement with the machine studies. The units differ by construction — responses to elicitation there, spontaneous citations in published text here — and our definition is strictly narrower: Dahl et al. [3] and Magesh et al. [8] count wrong-but-existent citations as hallucinations, which our closed-world method deliberately cannot see. Excluding the wrong-but-existent class makes our numerator a floor against a stricter standard; it also means the table's rates answer different questions, which is why we present them side by side and do not report a ratio.

4.2 Mechanism taxonomy

Because the series is closed, most errors match a local-damage signature. For every finding we evaluate deterministic rules — pure functions of the token and the series index, released with the code — in fixed priority order (Table 4). The labels are signatures, not causal attributions: 464 of 591 findings satisfy more than one rule, 115 exactly one, 12 none, and the table counts the first match in priority order. What the signatures establish is proximity: a small, structured

Table 4: Deterministic mechanism labels over all 591 findings: each label is a pure function of the token and the series index (Section 4.2). A finding may satisfy several rules; the first in priority order is counted.

Mechanism	DE	FR	IT	Total
Division substitution	92	169	11	272
Volume substitution	45	56	2	103
Extra digit in page	28	64	6	98
Year in page position	20	43	4	67
Dropped leading volume digit	16	14	0	30
Unlabelled residual	5	6	1	12
Doubled digit in page	4	4	1	9

edit maps the erroneous token onto an existing locus.

- **Division substitution** (272 findings, the dominant class): the cited page exists as a decision start under a sibling division of the same volume. Two criminal-division decisions print *BGE 148 I 356 E. 2.1, 39 E. 2.3.5* in the standard-of-review boilerplate — the continuation pair matches the criminal-division loci *148 IV 356* and *148 IV 39*, both of which exist.
- **Volume substitution** (103): same division and page as an exact start in a near or digit-edit volume. One court's footnote block cites *BGE 146 III 666* for the Pemetrexed doctrine while the neighbouring footnote cites *143 III 666* correctly; the deterministic repair recovers *143 III 666*.
- **Page-digit damage** (107): a doubled or extra digit, or transposed digits, in the page — *DTF 139 II 4040* for *139 II 404*.
- **Year for page** (67): *BGE 137 V 2010* for *137 V 210*; the year of decision intrudes into the page position.
- **Dropped leading volume digit** (30): *BGE 48 IV 137* for *148 IV 137* — provable only because the old volume lacks the cited division or page range; Section 5 bounds the sibling class that lands inside old volumes and escapes proof.

The rules propose repairs, and a repair asserts existence, not intent: 109 of 115 unique candidates exist on the court's resolver, but which locus the writer *meant* is a human judgment, made in this study only where the document itself answers it — in the appellate case of Section 3.3 the court states the intended citation (*BGE 149 III 210*), and the deterministic repair recovers exactly that token from *419 III 210*, a transposition. Within their limits the signatures still carry the structural point: these errors are small edits away from true citations, sometimes contradicted within the same document, where model fabrication invents whole plausible loci with no nearby ground truth — though the machine studies' definitions also cover misgrounding beyond invented loci (Table 3).

4.3 Propagation

89 distinct tokens account for 388 of the 591 findings; the most repeated, *BGE 142 I 455*, appears in 28 decisions. Errors travel inside reused doctrinal passages — the citation chains of standards of review, the footnote apparatus of recurring doctrine — and across court levels: *BGE 147 IIII 249* (division "IIII") recurs verbatim in three decisions from two court levels, each restating the same spousal-support standard. The distinct-locus view of Table 1 is the propagation-free reading of the same data: 292 distinct erroneous loci against 16,708 distinct cited loci, 1.75 % — and the reason single corrections at the source have outsized value.

Table 5: Sensitivity and strata. Every row divides errors by the matching edge universe from the same single-pass cluster structure as the headline; the 2026 row covers the partial year through the scan date.

View	Errors	Universe	Rate (ppm)
Headline	591	584,524	1,011
Excluding quotation-marked contexts	542	584,524	927
Federal courts	147	167,597	877
Cantonal courts	442	415,527	1,064
Other courts	2	1,400	1,429
Issued 2024	267	256,160	1,042
Issued 2025	257	237,454	1,082
Issued 2026	67	90,910	737

4.4 The remediation loop

Findings are not only counted. On 3 August 2026 the verified findings were reported to the issuing courts at federal and cantonal level, each letter carrying the verbatim context and, where evident, the intended citation. Within three days, seven court administrations replied: one confirmed the correction as already made, a second confirmed corrections under way, a third's official-collection unit committed to verifying and initiating corrections, and the remaining four acknowledged the findings and routed them to the responsible divisions. We deliberately do not name the courts: the point is the behaviour of the system, not a ranking, and the replies were courtesies, not statements for attribution. A daily automated audit now reports new findings the morning after publication. The claim this section supports is modest but, we believe, new: for a closed-world reporter series, the published record is not only auditable but *maintainable*. The facts of this section are author-reported from correspondence that is not public; they are dated, and the daily audit's outputs are released, but readers cannot verify the letters themselves.

5 Limitations

Reporter citations only. Decidability covers the BGE series; docket-form citations (*6B_1/2024*) are excluded because unpublished decisions make their universe open. The rate is therefore a statement about reporter citations, the most standardised class. **Lower bound.** A citation can be wrong yet existent (right page, wrong intended case); the closed-world method cannot see it, and the machine studies it is compared against do count that class (Table 3). The comparison is conservative in our disfavour. **Typos and asymmetric detectability.** Many findings are plainly typographic, and typos are not a confounder: a wrong locus in the official published text is a citation error under the same definition the LLM studies use, whatever its mechanical origin. But typo mechanics make detection asymmetric: a dropped digit whose result lands inside an old volume's page range escapes proof entirely — and worse, silently resolves, creating a semantically wrong graph edge. We measure the exposed surface: 1,646 resolved citations from the window (2,816 per million) target pre-1955 volumes. Restoring a dropped leading digit maps 278 of them into the modern volume's page range — 21 onto an exact decision start, the rest onto possible pin-cites. Those 278 citations, 47 % of the size of the finding set, are the unadjudicated risk pool of this one typo class; other wrong-but-existent channels (division, volume, page, semantic) are not bounded here at all. The pool, with per-row pre-screen flags and empty adjudication fields, is released. **Grammar coverage.** Citations written with the historical Ia/Ib sub-divisions are outside the extraction grammar (Section 3.2): 10,723 occurrences in the window, 1.4 % of the citation surface, all pointing at 1955–1994 volumes — the same old-material population as the typo asymmetry, and a further reason to read the headline as a floor. **Voice.** The quotation screen and single-reader pass of Section 3.3 found three party-attributed findings; reported speech without typographic markers remains possible. The sensitivity row excluding all marker-carrying findings moves the rate by about 8 %. No multi-annotator protocol was applied. **Source fidelity.** Our text layer can damage a correct citation; the window is born-digital, and among the 254 probed in-coverage tokens none resolved — but that verifies nonresolution of normalized tokens, not the official typography of every source occurrence, and pre-1954 tokens rest on the corpus alone. **One jurisdiction.** Switzerland's reporter structure enables the proof; the rate need not transfer, though the method does wherever a closed reporter series exists.

6 Data and code

Released with the corpus [5]: the scan, probe and grammar-count scripts; the findings file with per-finding context, mechanism labels, repair candidates and adjudication fields; the per-decision cluster structure; the complete deduplicated series index $(S_{v,d}, w_{v,d})$; the pre-1955 pool with pre-screen flags; the resolver probe results; the table generator; and a manifest with SHA-256 hashes of every artifact and of the exact database files the scan read. Every scan-derived number in this paper regenerates from these artifacts via a single script. Regeneration is formatting-level reproducibility; independently re-deriving the counts requires the corpus databases, whose content hashes the manifest pins. The correspondence facts of Section 4.4 are author-reported (Section 4.4).

References

[1] Fidan Badalova and Philipp Mayr. Detecting hallucinated and suspicious citations: What current tools can and cannot do. `https://arxiv.org/abs/2607.22693`, 2026. arXiv:2607.22693.

[2] Damien Charlotin. AI hallucination cases database. `https://www.damiencharlotin.com/hallucinations/`, 2026. Tracker of court and tribunal decisions confronting AI-generated hallucinations in filings; over 1,300 proceedings across 106 countries as of mid-2026. Accessed 2026-07-31.

[3] Matthew Dahl, Varun Magesh, Mirac Suzgun, and Daniel E. Ho. Large legal fictions: Profiling legal hallucinations in large language models. *Journal of Legal Analysis*, 16(1): 64–93, 2024. arXiv:2401.01301.

[4] Florian Geering and Jakob Merane. Swiss federal supreme court dataset (SCD), 2024. Version 2023-5.

[5] Jonas Hertner. OpenCaseLaw: An open multilingual corpus and citation graph for Swiss jurisprudence. `https://opencaselaw.ch`, 2026. Companion resource paper, in revision.

[6] Raizel Liebler and June Liebert. Something rotten in the state of legal citation: The life span of a United States Supreme Court citation containing an Internet link (1996–2010). *Yale Journal of Law and Technology*, 15:273–311, 2013.

[7] Patty Liu, Dominik Stammbach, and Peter Henderson. Who checks the citations? Benchmarking legal hallucination detection. `https://arxiv.org/abs/2606.21155`, 2026. arXiv:2606.21155.

[8] Varun Magesh, Faiz Surani, Matthew Dahl, Mirac Suzgun, Christopher D. Manning, and Daniel E. Ho. Hallucination-free? Assessing the reliability of leading AI legal research tools. *Journal of Empirical Legal Studies*, 22:216–242, 2025. doi: 10.1111/jels.12413. arXiv:2405.20362.

[9] Scott A. Mogull. Accuracy of cited facts in medical research articles: a review of study methodology and recalculation of quotation error rate. *PLOS ONE*, 12(9):e0184727, 2017.

[10] Volodymyr Ovcharov. Automatic construction of a legal citation graph from 100 million ukrainian court decisions: Large-scale extraction, topological analysis, and ontology-driven clustering. *arXiv preprint arXiv:2605.15362*, 2026. doi: 10.48550/arXiv.2605.15362.

[11] Princeton Center for Information Technology Policy. Can AI reduce burdens on courts by automatically verifying citations? `https://blog.citp.princeton.edu/2026/05/27/`, 2026. Blog essay, 27 May 2026.