

A Pilot Closing-Audit Pipeline for Legal RAG: Failure Analysis on Swiss Federal Law

Jonas Hertner¹

¹OpenCaseLaw — Independent researcher and maintainer, Zurich, Switzerland,
team@jonashertner.com <https://opencaselaw.ch>

April 29, 2026

Abstract

This is a *pilot failure-analysis note*, not a finalised verification system or a generalisable performance claim. We present a five-rail closing audit pipeline for retrieval-augmented legal LLMs — case-citation existence, statute-reference resolution, verbatim-quote source-matching, decision-date sanity, and proposition grounding by a separate-call LLM judge (same model family as the generator in the present experiment) — and stress-test it against Claude Sonnet 4.6 on 30 Swiss-federal-law questions in a *prior-only* condition (no retrieval). Four rails are deterministic (~ 3 ms per draft); the fifth is judge-mediated (~ 2 s, $\approx \$0.005$ per call). Three findings. (i) Citation existence is not the bottleneck in this setting: only 1 of 33 emitted citations is fabricated, and 5 of 6 drafts judged wrong by a same-family *c_{eli}* judge contain only valid citations; the misrepresented-holding class dominates. Comparable absolute hallucination rates reported by Dahl et al. 2024 and Magesh et al. 2025 were measured under different model/condition combinations and we do not claim a like-for-like comparison. (ii) On the 6-wrong-draft sample, the case-existence rail flags the one fabricated draft (0 % false-positive rate); adding the other deterministic rails preserves the wrong-draft flag rate (16.7 %) but raises FPR to 8.3 %; adding the grounding rail brings the wrong-draft flag rate to 50 % at FPR 25 %, but inspection shows neither of the two added flags is a clean semantic catch of the misrepresented holding — one fires on a secondary citation-attachment claim, the other on a fragmentary claim string. (iii) The 3 wrong drafts that escape the audit do so for mechanical reasons (claim-extraction heuristic skipped the citation on two; the judge call returned an error on one), not because the judge evaluates and concurs with the model’s mischaracterisation. We release the audit pipeline (MIT) and **Swiss Legal RAG Bench v0.2**, a 30-question pilot multilingual evaluation set (DE/FR/IT; benchmark packaging CC0). v1.0 will (a) implement more permissive claim extraction and judge-call retry, (b) add an end-to-end retrieval-augmented condition, (c) replace the same-family *c_{eli}* judge with an independent model, and (d) expand the benchmark to 150 questions with named expert co-authors and a multi-annotator protocol.

1 Introduction

Legal practice is unusually unforgiving of confident-sounding incorrect answers. A lawyer who cites a fabricated case, miscites a statute, or attaches a false holding to a real authority risks professional discipline and a wrong outcome for the client; in Switzerland, this is a violation of the *Sorgfaltspflicht* (duty of care) under the federal *Anwaltsgesetz* (BGFA, SR 935.61, art. 12). As large language models become a routine part of legal research, this asymmetry collides with a measured failure mode: 58–82 % of legal queries to general-purpose LLMs (Dahl et al., 2024) and 17–33 % of queries to commercial legal-RAG products (Magesh et al., 2025) produce a fabricated authority, miscited statute, or unsupported quotation.

Two recent developments shape the present work. First, Butler & Butler (2026) introduce a clean methodology for evaluating legal-RAG systems: decompose each error into three orthogonal binary signals (correctness c_{eli} , groundedness g_{eli} , retrieval accuracy r_{eli}) and read off the failure class from their joint distribution. Their methodology lets us ask *which* rail of a verification system addresses *which* error class. Second, the Swiss legal NLP literature has matured rapidly: SCALE (Rasiah et al., 2023) introduced a multi-task Swiss legal benchmark suite, the Niklaus group released anonymization tooling (Niklaus, Mamié et al., 2023) and re-identification risk measurement (Nyffenegger, Stürmer & Niklaus, 2024) for Swiss court rulings, SwiLTra-Bench (Niklaus et al., 2025) provided 180 k aligned Swiss legal translation pairs, and the SLDS dataset (Rolshoven, Rasiah et al., 2025) offered 20 k cross-lingual SFSC summarisation pairs. None of these target retrieval-augmented *generation* as the unit of evaluation; the present paper occupies that specific gap.

Two findings, in tension, motivate the rail decomposition. Our central empirical result is that on a closed Swiss-law corpus (969 k decisions, daily refresh; portions of which are plausibly present in the web-scale pretraining of the model under test) Claude Sonnet 4.6 fabricates only **3.0 %** of emitted citations. Comparable absolute hallucination rates have been reported by Magesh et al. (2025, 17–33 % on commercial legal-RAG products) and Dahl et al. (2024, 58–82 % on general LLMs); those measurements were taken under different model and condition combinations and we do not claim a like-for-like comparison. Citation existence is not the bottleneck in our setting. But the c_{eli} correctness rate is only 80 % (24/30): of the 6 wrong drafts, only 1 contains a fabricated citation; 3 are clean *real-citation*, *misrepresented-holding* errors (the model cites a real BGE correctly but omits its load-bearing precedential element); and 2 mix the misrepresentation pattern with a primary error of another class (an outdated statutory scale on q-018, a wrong-grounds enumeration on q-030). None of these five non-fabrication failures are detectable by per-citation existence checks; an LLM-judge grounding rail can in principle reach them, at meaningful false-positive cost, and in practice the rail is bottlenecked at our claim-extraction heuristic. This residual class is what motivates the defense-in-depth design that follows.

The conceptual move. A production verification system for legal RAG should not be a single black-box LLM judge over the final draft, nor a single per-citation retrieval check, but a **decomposition into independent rails, one per named error class**, each implemented with the cheapest mechanism that can detect that class. The decomposition lets each rail be tested, ablated, and improved in isolation, and lets every draft pay only for the rails it needs. We instantiate the design as a five-rail closing audit pipeline (§3) deployed in production at mcp.opencaselaw.ch.

Contributions.

1. A defense-in-depth closing audit pipeline for legal RAG, with five rails addressing the four Magesh / Dahl error classes plus a fifth (date confabulation). Latency ≤ 3 ms for the four deterministic rails; ~ 2 s, $\approx \$0.005$ per call for the LLM-judge rail.
2. **Swiss Legal RAG Bench v0.2:** a 30-question *pilot multilingual evaluation set* for Swiss federal law (DE / FR / IT; single-curator), modelled on Butler & Butler (2026) with one cross-lingual retrieval question (q-028) as a placeholder for a future cross-language suite. v1.0 (150 questions, named expert co-authors, multi-annotator protocol) is the planned gold-standard deliverable.
3. Per-rail ablation identifying two structural gaps in verification-first legal RAG: (a) the dominant residual error class is *real citation*, *misrepresented holding* (3 of 6 wrong drafts

cleanly, plus 2 that blend it with current-law / statutory-eligibility error, vs. 1 fabricated citation); (b) the audit’s claim-extraction heuristic skips the citation on a non-trivial fraction of drafts, blocking the grounding judge before it can evaluate. In our experiment the rail fires on the 2 wrong drafts where claim-extraction provided a candidate (though on inspection the trigger is a secondary citation-attachment issue rather than a clean semantic catch of the misrepresented holding) and is blocked on 3 of the 5 non-fabrication wrong drafts (claim-extraction skips $\times 2$, judge call error $\times 1$).

2 Related Work

Legal-LLM hallucination measurement. Dahl et al. (2024, *Large Legal Fictions*) profiled hallucination on six legal-task types in general-purpose LLMs, reporting fabricated-authority rates of 58–82 %. The follow-up Magesh et al. (2025, *Hallucination-Free?*) measured commercial legal-RAG tools (Lexis+ AI, Westlaw AI-Assisted Research, Ask Practical Law) at 17–33 %. Our paper accepts these as the problem statement; the five-rail pipeline of §3 is the treatment.

Legal-RAG benchmarks. Butler & Butler (2026, *Legal RAG Bench*) introduced the three-signal error decomposition (c_{eli} , g_{eli} , r_{eli}) we extend. Their benchmark covers Australian criminal law (4,876 passages from the Victorian Criminal Charge Book, 100 hand-crafted questions) and uses a closed corpus. We adapt the methodology to Swiss multilingual federal law, adding a cross-language retrieval signal because Switzerland is unusual among major jurisdictions in that the authoritative version of a leading case may be in a language other than the question’s (we acknowledge that v0.2 includes only one cross-lingual question, q-028; v1.0 expands this to a full cross-language retrieval suite). Closed-domain multiple-choice benchmarks (LegalBench-RAG, Pipitone & Hour Alami 2024; HousingQA / BarExamQA, Zheng et al. 2025; LegalBench, Guha et al. 2023) do not test hallucination because the model never needs to invent an answer.

Swiss legal NLP. Recent Swiss legal NLP has matured rapidly around the Niklaus / Stürmer / Stanford / UZH / Bern axis. Three lines are directly adjacent to ours. (i) *Privacy and anonymisation*: Niklaus, Mamié et al. (2023) released BERT-based models for entity-level anonymisation of Swiss FSCS rulings, improving the operational *Anom2* system used by the Federal Supreme Court; Nyffenegger, Stürmer & Niklaus (2024) systematically measured LLM-based re-identification risk on anonymised Swiss court rulings, finding the risk minimal for top models on *court decisions* (higher on Wikipedia entities). We discuss the implications for our corpus exposure in §6. (ii) *Multilingual benchmarks*: SwiLTra-Bench (Niklaus et al., 2025) is a 180k-pair Swiss legal translation benchmark across DE / FR / IT / RM / EN, introducing the SwiLTra-Judge specialised LLM evaluator with human-expert validation. SLDS (Rasiah et al., 2025) provides 20k SFSC rulings with cross-lingual headnote summarisations. We do not duplicate these tasks; our cross-lingual test (q-028) targets retrieval, not translation or summarisation. (iii) *Multi-task suites*: SCALE (Rasiah et al., 2023) covers citation extraction, court-view generation and summarisation; Swiss-Judgment-Prediction (Niklaus, Chalkidis & Stürmer, 2021) addresses outcome prediction; MultiLegalPile (Niklaus et al., 2023) provides a 689 GB multilingual legal pre-training corpus. None of these target retrieval-augmented *generation evaluation*; the present paper occupies that specific gap.

Verification methods. Self-RAG (Asai et al., 2023) and CRAG (Yan et al., 2024) propose general-domain self-verification or retrieval-correction; both treat verification as *pre-generation* (judge passage quality before generating). Our pipeline is *post-generation* verification: parse the generated draft, audit each emitted reference against the closed corpus, return a verdict per rail. The decomposition by named error class is, to our knowledge, novel for legal RAG;

it is analogous in spirit to the SwiLTra-Judge calibration approach (Niklaus et al., 2025) but addresses retrieval-grounded text generation rather than translation evaluation.

Positioning. To our knowledge, no prior work proposes a *deployed* mechanism that systematically reduces hallucination by decomposing the verification problem into independent error-class detectors. Existing RAG systems we are aware of treat verification as either a single LLM-judge pass over the final output (which inherits the same model’s biases) or per-citation retrieval validation (which catches existence errors but not propositional mis-grounding). The five-rail closing audit we propose in §3 is, to our knowledge, the first to address each error class with the cheapest mechanism that can detect it, with each rail independently testable and ablatable.

3 The Five-Rail Audit Pipeline

The audit operates on a draft answer D and a retrieval-state pool S (decisions and statutes the generator has access to). It returns a list of issues $\mathcal{I}(D, S)$, each labelled with one of five categories (Figure 1). An OK verdict (no issues) means the audit found nothing; $\neg$ OK means at least one rail fired. The five rails operate *in parallel* on the same draft.




Figure 1: Five-rail closing audit pipeline. Four deterministic rails address citation-, statute-, quotation-, and date-level fabrications; an opt-in LLM-judge rail addresses proposition-grounding errors. Each rail is independently testable and ablatable; the configuration adopted by a deployment is a per-rail subset.

Rail 1 — case-citation existence. A regex parses the draft for Swiss case-citation forms (BGE / ATF / DTF, BGer / TF dockets, BVGer / BStGer / BPatGer dockets, MKGE numerical references). Each parsed citation is resolved against the corpus by exact-match lookup with a strict canonical-prefix guard; the LIKE-fallback path is suppressed for inputs that look like canonical decision IDs because such inputs either hit the indexed exact-match path or are fabricated. Latency: 2–5 ms per draft regardless of citation count. Catches: invented BGE numbers, mis-typed docket numbers, citations to retracted decisions. Optional pinpoint check: if the citation includes E.*n.m* or consid.*n.m*, the pinpoint must match either the structured-extraction sidecar or the body-text **regeste** of the cited decision.

Rail 2 — statute-reference resolution. A second regex parses for Art.*N* LAW forms (with German / French / Italian abbreviation synonyms) and resolves each against the Fedlex SPARQL mirror cached locally as **statutes.db**. The resolution is two-step — abbreviation → SR number → article text — and is cached at process scope. Latency: 1–3 ms per draft. Catches: fabricated law abbreviations, wrong article numbers, references to repealed articles. Connector words (“und”, “oder”, “et”, “ou”) are case-sensitively filtered to prevent the regex eating them as fake law codes after Abs.*N* structures.

Rail 3 — verbatim-quote source-matching. Quoted substrings of ≥ 30 characters in the draft are extracted (German „...“, French « ... », Italian « ... », English curly “...”, ASCII “...”) and required to appear verbatim — after whitespace + quote-mark normalisation — in the

source pool S . The pool comprises the **regeste** text plus first 8k characters of each cited decision, plus the article text of each cited statute. Latency: ~ 2 ms. Catches: hallucinated quotations, paraphrases inside quotation marks.

Rail 4 — decision-date sanity. For each verified case citation, parse adjacent vom / du / del DD.MM.YYYY constructs within 60 characters of the citation span and compare to the stored **decision_date**. Catches: model emits a real BGE but pairs it with a wrong date, suggesting confusion between two cases with similar dockets. Latency: < 1 ms.

Rail 5 — proposition grounding (opt-in). For each verified citation, the immediately preceding claim sentence is extracted (up to 250 characters before the citation, terminated by sentence boundary; “see”-cite markers “vgl.” / “siehe” / “cf.” / “cfr.” suppress the rail). All (claim, source) pairs go in a *single batched* prompt to a separate Sonnet judge call (same model family as the generator; see §6) that returns a JSON array of verdicts {**yes**, **partial**, **no**, **contradicts**, **unrelated**}. Verdicts in {**no**, **contradicts**, **unrelated**} become issues. One Sonnet call per draft regardless of citation count. Latency: 1.5–2.5s. Cost: $\approx$ \$0.005 per call. Targets the Butler / Butler “reasoning error” class (real citation + misrepresented holding); empirical catch rate on this class is qualified in §5 (the rail in our experiment fires on secondary attachment issues rather than cleanly on the underlying mischaracterisation).

Output canonicalisation. When all rails pass, the audit emits a **linked_text** field with every parsed citation replaced by its corpus-canonical form wrapped in a Markdown link to the source URL on mcp.opencaselaw.ch/entscheid/<id>. Drafts arrive with messy citation spacing or case and leave publication-ready.

4 Swiss Legal RAG Bench v0.2

We release Swiss Legal RAG Bench v0.2 alongside this paper as a *pilot multilingual evaluation set*. Version 0.2 contains 30 questions *author-curated by the maintainer of opencaselaw.ch (the corpus mirror underlying the audit pipeline)*, drawn from the major substantive areas of Swiss federal law (Schuldrecht, Sachenrecht, Familienrecht, Erbrecht, Strafrecht, Strafprozess, Verfassungsrecht, Verwaltungsrecht, Mietrecht, Arbeitsrecht, Sozialversicherungsrecht, Verkehrsrecht, Aktienrecht, SchKG, Datenschutz, Verjährung, Wirtschaftsfreiheit; 17 distinct areas). The language distribution is DE = 18, FR = 8, IT = 4 (60 / 27 / 13 %). v0.2 is explicitly framed as a single-curator seed; v1.0 will adopt a multi-annotator protocol with named expert co-authors from Swiss law faculties (see §6).

Question construction. Each question is paired with a concise reference answer (1–4 sentences in the question’s language) and annotated supporting evidence (statute SR-number + article number + language; leading-case **decision_id**(s) where applicable). Of the 30 questions, 16 (53 %) are anchored on a verified leading-case BGE; 14 (47 %) carry the **doctrine_from_bge** claim type, in which the expected answer derives a doctrinal proposition from a seed BGE rather than restating its holding directly. The two annotations are not disjoint: every **doctrine_from_bge** question is also anchored on a leading-case BGE. The split exists to let the ablation surface where each rail contributes — case-existence on the seed citations, the grounding judge on the doctrinal proposition.

Corpus verification of evidence. Every **decision_id** cited as evidence was resolved against the live corpus (mcp.opencaselaw.ch/api/decisions/<id>) before the benchmark was published; every BGE returned a 200 response with non-zero citation count and an on-topic **regeste**.

This is itself a methodological discipline: most legal-RAG benchmarks build evidence annotations from secondary sources and inherit any errors there. We commit to corpus-first construction, with the audit trail published in the supplementary material.

Cross-lingual extension. Switzerland is unusual among major civil-law jurisdictions in that the authoritative version of a leading case may be in a language other than the question’s. Question q-028 (IT) anchors on DTF 126 III 113 (which has *regesti* in DE / FR / IT) to test the cross-language retrieval claim. v1.0 expands this dimension to a full cross-language retrieval suite.

Versioning. v0.2 is a seed sufficient to characterise the audit pipeline in this paper. **v1.0** is targeted at 150 questions with named expert co-authors from Swiss law faculties and a multi-annotator protocol (inter-annotator agreement reported, third-party adjudication for disagreements). We explicitly invite contributors for v1.0.

5 Experiments

Setup. We use Claude Sonnet 4.6 (released 2026-02-17) as both the generator and the judge for c_{eli} , g_{eli} , and the in-pipeline grounding rail. Using one model family on both sides is a known limitation (§6); v1.0 will report inter-judge agreement against an independent model. All 30 questions in v0.2 are run in a *prior-only* condition (no retrieval): the generator sees only the question text and must rely on training knowledge. We chose the prior-only condition as an audit-pipeline *stress test*, not as an end-to-end RAG evaluation: it maximises the density of recoverable errors, and the no-retrieval setup is methodologically aligned with Dahl et al. 2024’s general-LLM hallucination measurement (we do not claim it replicates Dahl’s substantive coverage of US legal tasks; our questions are Swiss federal law, single-curator). A retrieval-augmented condition runs in production at mcp.opencaselaw.ch (BM25 / RRF / Haiku rerank, see §6); end-to-end RAG evaluation against v1.0 is forthcoming work.

For each question we collect: the prior-only draft, the per-category issue counts from `attest_response` with all five rails on, the c_{eli} judge verdict against the v0.2 reference answer, and the g_{eli} judge verdict against the empty source pool. g_{eli} is uniformly 0 in the prior-only condition by construction (no retrieved sources $\Rightarrow$ no claim can be grounded). We use the c_{eli} verdict as ground truth for the per-rail ablation.

Per question we issue exactly four Sonnet calls (one generation, one in-pipeline grounding judge, one c_{eli} judge, one g_{eli} judge) plus one deterministic-rails pass; the 6-row ablation in Table 1 is computed in post-processing by intersecting each draft’s per-rail fire flags with each configuration’s enabled rail set. Total: 120 Sonnet calls across 30 drafts; 106s wall-clock with 4-way parallelism; <\$2 in API charges.

Citation-level fabrication rate. Across 30 prior-only drafts, Claude Sonnet 4.6 emits 33 case citations of which 32 resolve in the corpus (citation-level fabrication: **3.0%**). More informatively, of the 6 drafts judged wrong by the c_{eli} judge, 5 contain only valid citations: *citation existence is not the bottleneck in this setting*. Magesh et al. (2025) reported 17–33% on commercial legal-RAG tools and Dahl et al. (2024) reported 58–82% on general LLMs; those measurements were taken under different model families, query distributions, and tool conditions than ours, so the absolute numbers should not be read as a like-for-like comparison. Two non-exclusive hypotheses, neither tested here: (i) Swiss BGE references appear at scale on the open web — including via our own published Hugging Face dataset — and the model may have internalised the citation conventions; (ii) the prior-only condition asks short, well-bounded questions where the model can refuse, hedge, or ground in widely-known leading cases, whereas Magesh’s tools were exercised on more open-ended litigation queries.

Per-rail ablation. Of the 30 prior-only drafts, 24 are judged correct ($c_{eli} = 1$) and 6 wrong ($c_{eli} = 0$) by the second-pass Sonnet judge (same model family as the generator; see §6). Table 1 reports the rate at which each cumulative configuration *flags* drafts in the wrong-draft subset (**wrong-draft flag rate**, “WFR”) and the correct-draft subset (**false-positive rate**, FPR). We deliberately avoid calling WFR a “true-positive rate” or “catch rate” because, as the inspection in this section makes clear, several flags happen to land on wrong drafts without semantically detecting the underlying primary error.

Table 1: Per-rail-configuration flag rates, v0.2 prior-only condition. $n = 30$ drafts; $c_{eli} = 0$ for 6 (wrong) and $c_{eli} = 1$ for 24 (correct) per the second-pass Sonnet judge. WFR (wrong-draft flag rate) = $P(\text{any enabled rail fires} | c_{eli} = 0)$, computed on $n = 6$; FPR = $P(\text{any enabled rail fires} | c_{eli} = 1)$, computed on $n = 24$. Net = WFR – FPR. WFR counts *any* rail firing on a wrong draft and does *not* certify that the flag identifies the draft’s primary error; the post-table inspection shows neither of the two grounding-rail flags is a clean semantic catch of the misrepresented holding. 95 % Wilson CIs on WFR are wide because of the small wrong-draft denominator ($3/6 \rightarrow [18.8, 81.2]$).

Config	Rails enabled	WFR	WFR 95% CI	FPR	Net
C_0	$\emptyset$ (no rails)	0/6 = 0.0 %	[0.0, 39.0]	0.0 %	+0.0 pp
C_1	+ case	1/6 = 16.7 %	[3.0, 56.4]	0.0 %	+16.7 pp
C_2	+ statute	1/6 = 16.7 %	[3.0, 56.4]	4.2 %	+12.5 pp
C_3	+ quote	1/6 = 16.7 %	[3.0, 56.4]	8.3 %	+8.3 pp
C_4	+ date (4 deterministic)	1/6 = 16.7 %	[3.0, 56.4]	8.3 %	+8.3 pp
C_5	+ grounding (full 5-rail)	3/6 = 50.0 %	[18.8, 81.2]	25.0 %	+25.0 pp

Three findings emerge. First, the deterministic rails behave very differently from one another: only the case-existence rail (configuration C_1) holds the perfect-precision floor (0 % FPR), and it does so by catching a single fabrication, q-017’s invented DTF 117 II 198, which Sonnet uses to anchor a substantively flawed art. 264 CO answer (the draft also omits cpv. 3 and mischaracterises the unjustified-refusal consequences — both correctness failures the case-existence rail is not designed to detect). Adding the statute, quote, and date rails (C_2 – C_4) preserves TPR but degrades precision: FPR climbs to 8.3 % by C_4 as the rails flag valid statute references and quotations as anomalies on correctly-answered drafts. Second, the grounding rail (C_5) is responsible for the entire jump in TPR over C_4 : it fires on the only two wrong drafts (q-018, q-030) where claim-extraction supplied the judge with a candidate. Inspection qualifies the credit: on q-018 the rail flags a secondary citation-attachment claim (“BGE 135 V 297” attached to a generic concretisation claim) while the draft’s primary correctness failures — omission of *Eingliederungsmassnahmen*, the six-month waiting period, and the 2022 linear scale — go undetected; on q-030 the extracted claim is the fragment “Entscheidend ist nach”, and the judge returns **unrelated** because the claim string is too short to evaluate. Both flags happen to land on wrong drafts, but neither is a clean semantic catch of a misrepresented-holding error. The 25 % FPR is the cost of that recall: six correct drafts trigger the rail because the second-pass judge disagrees with the draft’s interpretation of a real cited authority. Third, the three wrong drafts that all rails miss (q-021, q-025, q-026) are missed for a structural reason different from the explanation we originally hypothesised.

The residual error class. The three missed drafts (q-021, q-025, q-026) share a substantive pattern — Sonnet cites a real BGE (BGE 142 III 433, BGE 144 III 481, BGE 130 V 343) correctly but *mischaracterises* the holding, omitting the load-bearing precedential element — which we call the *real-citation / misrepresented-holding* class. However, the proximate reason these drafts escape the audit is *not* that the grounding judge evaluates them and concurs with Sonnet’s mis-

characterisation: it is that the grounding rail never gets to evaluate them. Per `grounding_meta` in the artifact: q-021 and q-026 trip the `skipped_no_claim` branch (the heuristic that extracts the claim sentence preceding a citation found no candidate within its 250-character window), and q-025 returns a `judge_unavailable` error (a transient API failure during the experimental run). The residual error class is therefore really two failures stacked on top of each other: an audit-pipeline gap (claim-extraction heuristic, judge-call reliability) prevents the verifier from *evaluating* the drafts that contain the most subtle error type, and the substantive error type itself is the hardest one even given evaluation. Closing the audit gap is mechanical (more permissive claim extraction, retry-on-error) and is the obvious next deployment step; testing whether the grounding judge actually catches misrepresented-holding errors when it gets to evaluate them requires v1.0’s larger denominator and an independent judge model.

Cost / latency. Across the 30 drafts, the four deterministic rails together account for ~ 3 ms per draft (median 2.7 ms; 95 % under 6 ms), measured by wrapping the rail dispatch in `time.perf_counter_ns()` on the production deployment. The grounding rail takes 1.5–2.5 s wall-clock and $\approx \$0.005$ per draft via the Anthropic API. Total experimental cost (generation + audit + judges): $< \$2$.

6 Discussion

When the deterministic rails miss. The case-citation rail catches *fabricated* references, not *misused* ones. For systems where the underlying model has high citation-level fidelity (as Sonnet 4.6 has on Swiss law: 3.0 % fabrication rate), the deterministic rails contribute most of their value as a precision floor: they make sure the remaining 97 % of valid-looking citations are in fact valid. Of the four deterministic rails, only C_1 (case-existence) holds 0 % FPR in our experiment; adding statute, quote, and date raises FPR to 8.3 %. A production deployment can run all four cheaply, but the precision-floor argument for running them on every draft applies cleanly only to C_1 as currently implemented; the other three need to be tightened or made user-suppressible before they fire as freely.

When the judge rail miscalls. The 25 % FPR of the grounding rail in our experiment plausibly overstates the production rate, but not for the reason we initially supposed: the rail already fetches the cited regeste / source text for verified citations, so the judge’s access to source material does not change with retrieval. What changes with retrieval is the *generator*: a draft built on retrieved passages tends to anchor each claim with a pinpoint that maps cleanly to the source, leaving the judge less room to score `partial` or `unrelated`. This is a hypothesis, not a result of the present paper. The fundamental tradeoff remains in either condition: a stronger judge catches more reasoning errors but raises more disagreement-as-error verdicts; calibrating it — when to defer, when to override — is open work.

Generalisation. Each rail is corpus-agnostic in design. The case-existence rail needs an enumerable decision corpus; the statute rail needs a statute mirror; the quote rail needs the cited decisions’ verbatim text. None of these are Swiss-specific. An English-law adaptation (e.g. England & Wales) would require: (i) regex patterns for English neutral citations and law-report references; (ii) a statute mirror for primary legislation (e.g. `legislation.gov.uk`-derived); (iii) decision full-text for the quote rail. Open or low-friction sources exist for much of this — BAILII for judgments, `legislation.gov.uk` for primary legislation — but ICLR / commercial law-report text sits behind licensing, and corpus-construction licensing is jurisdiction-specific in ways we do not attempt to inventory here. We expect the catch-rate / FPR profile to be model-dependent: a weaker generator with higher citation-level fabrication rate will see higher rail TPR.

Privacy and re-identification. The corpus we redistribute (969k Swiss court decisions; packaging and added metadata under CC0; underlying texts are official published court decisions of the issuing courts) consists exclusively of decisions *as published by the issuing court* or via the official Swiss judicial publication portal `entscheidsuche.ch`. We perform no de-anonymisation, no re-identification, and no enrichment with non-public data. Nyffenegger, Stürmer & Niklaus (2024) measured LLM-based re-identification risk on anonymised Swiss court rulings and found it minimal for top frontier models on *court decisions* specifically (substantially higher on Wikipedia entities); we adopt their finding and discuss it explicitly in the dataset card. Our contribution to privacy is one-direction: we expose the already-public anonymised corpus through a more uniform interface, which does not create re-identification capacity beyond what was already feasible by direct scraping of the source portals. We do *not* mirror the in-court (pre-anonymisation) corpus; that work — including the operational *Anom2* system — belongs to the Niklaus, Mamić et al. (2023) line and the Federal Supreme Court itself.

Limitations.

- **Pilot scope.** This paper is a pilot failure-analysis note, not a finalised verification system or a generalisable performance claim. Several of the audit-pipeline behaviours we report are visibly fixable (claim-extraction skip, judge-call retry; see below), and we explicitly chose to publish the as-instrumented numbers *before* fixing them so that the failure modes are documented for the community rather than papered over in code. v1.0 will rerun after the fixes.
- **Mechanical audit-pipeline gaps.** Three of the six wrong drafts (q-021, q-025, q-026) are missed by all rails for mechanical reasons: the claim-extraction heuristic looks for a sentence-bounded candidate within 250 characters before each citation and finds none on q-021/q-026; the grounding judge call on q-025 returned a `judge_unavailable` error (no retry was attempted). Both gaps are addressable in code; we report the as-instrumented numbers here so the failure modes are documented faithfully against the artefact pinned in the Data and Code Availability section, and defer the fixed-code rerun to v1.0 to avoid moving the experimental target while also drafting the paper. We do *not* claim the post-fix numbers in advance of running them.
- **Benchmark size.** v0.2 is a 30-question seed. With only 6 wrong drafts, the per-configuration WFR estimates are imprecise: 95 % Wilson confidence intervals on a 50 % point estimate ($n = 6$) are [18.8%, 81.2%]. We report exact ratios in Table 1 so the reader can judge directly. v1.0 (150 questions, named expert co-authors from Swiss law faculties; see §4) is the planned next deliverable.
- **Single-curator benchmark.** v0.2’s questions and reference answers were authored by the maintainer of `opencaselaw.ch`, a Swiss-trained author with practical legal-research experience but without the multi-annotator protocol that a published gold-standard benchmark requires. Selection bias is the obvious risk: a curator who built the underlying corpus may unintentionally choose questions whose evidence the corpus answers well. We mitigate by anchoring every evidence `decision_id` on the live corpus before publication and by publishing the evidence-resolution audit trail; we cannot eliminate the bias, and v1.0 addresses it directly with named expert co-authors and an independent annotation protocol.
- **Same model family, generator and judge.** Both roles are filled by Claude Sonnet 4.6 in our experiment. This concentrates failure modes (a wrong draft and a wrong-but-self-consistent judge will agree). v1.0 will adopt a SwiLTra-Judge-style independent calibrated judge (Niklaus et al., 2025) with human-expert validation, and report inter-judge agreement.

- **Ground-truth circularity.** We use the c_{eli} verdict (also produced by Sonnet 4.6) as the ground-truth label for which drafts are “wrong” in the per-rail ablation (Table 1). If Sonnet has a systematic bias on a given proposition, that bias affects both the draft and its c_{eli} judgment in the same direction, and a draft can be scored “correct” because the same misconception underlies both. We partially mitigate this by anchoring each question’s reference answer on a corpus-verified BGE before generation, so the c_{eli} judge has an explicit external anchor to compare against; we do not claim to have eliminated the circularity. v1.0’s independent judge addresses this directly.
- **Prior-only condition.** Our reported ablation runs the generator without retrieval, chosen as a stress-test of the audit pipeline (it maximises the density of recoverable errors), not as an end-to-end RAG evaluation. Butler & Butler (2026) report that retrieval is the primary determinant of legal-RAG performance; we do not contradict them, and an end-to-end retrieval-augmented evaluation against v1.0 is forthcoming. The audit pipeline is run on every production query at mcp.opencaselaw.ch (over a BM25 / RRF / Haikurerank retriever), but production-side audit catch-rates are not the subject of the present paper.
- **Single-run experiment.** The numbers in Table 1 reflect one run; we do not bootstrap. Run-to-run variance at `temperature=0` is small but non-zero.

7 Conclusion

This pilot proposes a defense-in-depth closing audit pipeline for legal RAG, decomposing verification into five rails with bounded latency and cost, and reports a 30-question stress test of an early instrumentation of that pipeline. We release the audit pipeline (MIT) and Swiss Legal RAG Bench v0.2 (CC0 packaging), the latter as a 30-question pilot multilingual evaluation set for Swiss federal law. In our prior-only stress test, the case-existence rail catches the one fabricated citation among 6 wrong drafts at 0% FPR; the second-pass grounding rail fires on the 2 wrong drafts where claim-extraction supplied a candidate (though on inspection both flags target secondary attachment issues, not the primary misrepresentation); the remaining 3 wrong drafts escape the audit because the claim-extraction heuristic skipped the citation ($\times 2$) or the judge call returned an error ($\times 1$) before any verdict was reached. Closing those audit-pipeline gaps, calibrating the grounding judge against an independent model on a larger-denominator v1.0, and running an end-to-end retrieval-augmented evaluation are the next steps.

The audit pipeline is live at mcp.opencaselaw.ch, deployed on the entire published Swiss federal and cantonal court corpus (969,000+ decisions, daily refresh). Code under MIT; benchmark packaging under CC0; underlying decision texts are official published court decisions of their issuing courts (with the collation, deduplication, and added metadata licensed CC0). We invite collaborators for v1.0.

Data and Code Availability

Reproducibility pin (2026-04-28). The numbers in §5 were produced from these exact artefacts:

- Audit pipeline source (MIT), pinned to the experiment commit:
https://github.com/jonashertner/caselaw-repo-1/blob/4553379/mcp_server.py
- Run harness (the script that generated 30 prior-only drafts and audited them):
https://github.com/jonashertner/caselaw-repo-1/blob/4553379/docs/paper/v2/experiments/01_run_ablation.py

- Table-rendering script (computes `ablation_table.json/.md` from the per-trial results in post-processing):
https://github.com/jonashertner/caselaw-repo-1/blob/4553379/docs/paper/v2/experiments/02_compute_table.py
- Per-trial ablation data, pinned to the experiment commit:
https://github.com/jonashertner/caselaw-repo-1/blob/4553379/docs/paper/v2/experiments/ablation_results.json
SHA-256: 79dcc442...039c3b08
- Derived ablation table (the JSON behind Table 1):
https://github.com/jonashertner/caselaw-repo-1/blob/4553379/docs/paper/v2/experiments/ablation_table.json
- Benchmark v0.2 (CC0): <https://huggingface.co/datasets/voilaj/swiss-legal-rag-bench>
SHA-256 of `questions.jsonl`: b299db37...f21ec449
- Corpus-verification audit trail (the evidence-resolution log produced when each question’s BGE evidence was checked against the live corpus before publication):
https://github.com/jonashertner/caselaw-repo-1/blob/4553379/docs/paper/v2/03c_verification_results.md
- Underlying Swiss-caselaw corpus (packaging and added metadata under CC0; underlying texts are official published court decisions of the issuing courts): <https://huggingface.co/datasets/voilaj/swiss-caselaw>
Live MCP corpus size at experiment time: 969,738 decisions. The HF dataset is updated daily; for byte-exact reproduction, point the harness at the live MCP endpoint with the commit-pinned audit code, which records each resolved citation’s `decision_id` so the experiment is snapshot-aware against the corpus state at run-time.
- Live MCP endpoint (for re-execution against current corpus): <https://mcp.opencaselaw.ch>

Generator and judge: `claude-sonnet-4-6` (Anthropic), `temperature=0`. Each trial in §5 can be re-executed by pinning to the commits + SHAs above and pointing the harness at the live MCP endpoint.

No competing financial interests declared. This work was funded by the author and the volunteer [OpenCaseLaw contributor community](#). The author maintains the project as a public-interest infrastructure; no commercial product is sold.

References

- [1] A. Asai, Z. Wu, Y. Wang, A. Sil, H. Hajishirzi. Self-RAG: Learning to retrieve, generate, and critique through self-reflection. *arXiv:2310.11511*, 2023.
- [2] A.-R. Butler, U. Butler. Legal RAG Bench: an end-to-end benchmark for legal RAG. Isaacus, March 2026. *arXiv:2603.01710*.
- [3] M. Dahl, V. Magesh, M. Suzgun, D. E. Ho. Large Legal Fictions: Profiling Legal Hallucinations in Large Language Models. *Journal of Legal Analysis*, 16(1), 64–93, 2024. *arXiv:2401.01301*.
- [4] M. Geering, J. Merane. Swiss Federal Supreme Court Dataset, 2024. Hugging Face datasets repository, `rcds/swiss_court_view_generation`.

- [5] N. Guha, J. Nyarko, D. E. Ho, et al. LegalBench: A collaboratively built benchmark for measuring legal reasoning in large language models, 2023.
- [6] V. Magesh, F. Surani, M. Dahl, M. Suzgun, C. D. Manning, D. E. Ho. Hallucination-Free? Assessing the Reliability of Leading AI Legal Research Tools. *Journal of Empirical Legal Studies*, 22(2):216–242, 2025. DOI: 10.1111/jels.12413. Working-paper version: Stanford RegLab, 2024 (arXiv:2405.20362).
- [7] J. Niklaus, I. Chalkidis, M. Stürmer. Swiss-Judgment-Prediction: A Multilingual Legal Judgment Prediction Benchmark, 2021.
- [8] J. Niklaus, V. Matoshi, M. Stürmer, I. Chalkidis, D. E. Ho. MultiLegalPile: A 689 GB Multilingual Legal Corpus, 2023.
- [9] J. Niklaus, R. Mamié, M. Stürmer, D. Brunner, M. Gygli. Automatic Anonymization of Swiss Federal Supreme Court Rulings. In *Proceedings of the Natural Legal Language Processing Workshop 2023*, ACL.
- [10] A. Nyffenegger, M. Stürmer, J. Niklaus. Anonymity at Risk? Assessing Re-Identification Capabilities of Large Language Models in Court Decisions. In *Findings of NAACL 2024*, pages 2433–2462.
- [11] J. Niklaus, J. Merane, L. Nenadic, S. Ahmadi, Y. Gao, et al. SwiLTra-Bench: The Swiss Legal Translation Benchmark. In *Proceedings of ACL 2025*, pages 14894–14916.
- [12] L. Rolshoven, V. Rasiah, S. Brügger Bose, S. Hostettler, L. Burkhalter, M. Stürmer, J. Niklaus. Unlocking Legal Knowledge: A Multilingual Dataset for Judicial Summarization in Switzerland. In *Findings of EMNLP 2025*, pages 15382–15411.
- [13] N. Pipitone, G. Houir Alami. LegalBench-RAG: A benchmark for retrieval-augmented generation in the legal domain. arXiv:2408.10343, 2024.
- [14] V. Rasiah, R. Stern, V. Matoshi, M. Stürmer, I. Chalkidis, D. E. Ho, J. Niklaus. SCALE: Scaling up the Complexity for Advanced Language Model Evaluation, 2023.
- [15] S. Yan, J. Gu, Y. Zhu, Z. Ling. Corrective Retrieval Augmented Generation. arXiv:2401.15884, 2024.
- [16] L. Zheng, N. Guha, J. Arifov, S. Zhang, M. Skreta, C. D. Manning, P. Henderson, D. E. Ho. A reasoning-focused legal retrieval benchmark. In *Proceedings of the Symposium on Computer Science and Law*, ACM, 2025.